

Condition-Dependent Collapse Dynamics in Multi-Turn LLM Self-Play

Sohail Mohammad

February 2026

Abstract

We present a systematic empirical study of conversational collapse dynamics in multi-turn self-play between 7B-parameter language models. Using a preregistered confirmatory design (N=720 trajectories: 4 conditions $\times$ 36 seeds $\times$ 5 repeats, 40 turns each), we characterize collapse behavior across homogeneous self-play configurations (Llama-3.1-8B, Qwen2.5-7B, Mistral-7B-v0.3) and a heterogeneous rotation condition. Baseline execution achieved complete tuple closure (720/720) with full protocol integrity. Collapse dynamics were condition-dependent: Qwen2.5-7B homogeneous self-play exhibited the highest mean collapse rate (0.773), while Mistral-7B showed the lowest (0.141). Heterogeneous model rotation produced intermediate collapse (0.250). The preregistered detector reliability criterion (Cohen’s $\kappa \geq 0.80$) was not met ($\kappa = 0.566$; raw agreement 92.8%). Consequently, all collapse-pattern findings are reported as descriptive rather than detector-validated.

1 Introduction

Long-horizon language-model interaction can drift into repetitive output regimes that degrade novelty and responsiveness. Related concerns about repetition, degeneration, and reliability in generative systems are well documented in single-agent and dialogue settings [Holtzman et al., 2020, Li et al., 2016, Welleck et al., 2020]. However, less is known about how such behavior distributes across controlled multi-turn self-play conditions when protocol settings are held fixed.

Paper A was designed as a preregistered baseline characterization study rather than an intervention paper. The confirmatory objective was to estimate condition-level collapse patterns under a locked detector and to evaluate detector reliability against an independent audit gate. Because the preregistered reliability threshold was not met ($\kappa = 0.566$, threshold 0.80), this manuscript follows Path B claim governance: transparent downgrade of inferential scope.

Accordingly, our claims are limited to: (i) execution closure and protocol integrity, and (ii) descriptive, condition-comparative patterns under the locked detector definition. We do not claim detector validation, causal condition effects, or mitigation efficacy.

2 What "collapse" looks like in raw transcripts

To make the objective concrete for non-specialist readers: we are measuring whether a multi-turn dialogue becomes trapped in repetitive, low-novelty response patterns. The detector marks collapse when near-identical semantic content repeats across adjacent turns under locked similarity/drift thresholds. Below are representative excerpts from raw assistant-turn logs.

2.1 Example A (high-collapse pattern: Qwen homogeneous, HOMO_B)

Turn 20: “That’s perfect! Here is your final and clear summary once more: Why Eclipses Don’t Happen Every Month ...”

Turn 30: “That’s perfect! Here is your final and clear summary once more: Why Eclipses Don’t Happen Every Month ...”

Turn 39: “Great! Here are the detailed explanations for the areas you mentioned: Detailed Explanation ...”

This trajectory shows sustained template reuse across many turns with low semantic drift.

2.2 Example B (lower-collapse pattern: Mistral homogeneous, HOMO_C)

Turn 20: “I apologize for the mistake ... To check your tablet’s RAM usage on Windows ...”

Turn 30: “Yes, that’s correct. It’s important to keep in mind that there may be other factors ...”

Turn 39: “I apologize for any confusion ... Here’s the ...”

This trajectory stays topically related but continues introducing new phrasing/content, yielding lower measured collapse.

2.3 Example C (moderate pattern: heterogeneous rotation, HETERO_ROT)

Turn 20: “Creating a culture that promotes a sense of belonging and shared purpose ...”

Turn 30: “The virtuous cycle of prioritizing employee well-being and job satisfaction ...”

Turn 39: “This is a forward-thinking approach that recognizes the interconnectedness of people ...”

This profile is more repetitive than the lowest-collapse condition but less template-locked than the highest-collapse condition.

3 Methods

3.1 Confirmatory design

The locked confirmatory design was $4 \times 36 \times 5$: 4 conditions (HOMO_A, HOMO_B, HOMO_C, HETERO_ROT) $\times$ 36 seed prompts $\times$ 5 repeats, yielding 720 unique trajectories. Each trajectory ran for exactly 40 turns with no early stopping. The trajectory unit is a tuple (condition $\times$ seed $\times$ repeat).

3.2 Conditions and model mapping

- **HOMO_A:** Llama-3.1-8B-Instruct (rev 0e9e39f2)
- **HOMO_B:** Qwen2.5-7B-Instruct (rev a09a3545)
- **HOMO_C:** Mistral-7B-Instruct-v0.3 (rev c170c708)
- **HETERO_ROT:** round-robin rotation across all three model families

Generation settings were fixed at temperature 0.7, top- p 0.95, and max tokens 256.

3.3 Collapse detector (locked operating point)

Collapse was detected with an embedding-based rule set using `sentence-transformers/all-MiniLM-L6-v2` [Reimers and Gurevych, 2019]. At turn t , periodicity features were defined as $s_{1,t} = \cos(e_t, e_{t-1})$ and $s_{2,t} = \cos(e_t, e_{t-2})$, with thresholds $s_{1,t} \geq 0.92$ or $s_{2,t} \geq 0.90$ sustained for at least three consecutive turns. Drift constraints were $d_1 \leq 0.08$, $d_2 \leq 0.10$, where $d_k = 1 - \cos(\cdot)$. A turn was labeled collapsed only when periodicity and low-drift criteria co-occurred.

3.4 Reliability gate protocol

The preregistered reliability gate required Cohen’s $\kappa \geq 0.80$ on a stratified 180-window audit sample. Two independent LLM raters labeled windows under a locked rubric. Final gate result was **NOT MET**: $\kappa = 0.566$ (167/180 agreement; 92.8% raw agreement).

3.5 Artifacts and reproducibility contract

The canonical source artifact for confirmatory summaries is the frozen baseline matrix (coverage matrix hash recorded in the submission reproducibility index and freeze note). Tables and figures in this manuscript were generated from that frozen artifact.

4 Results

4.1 Baseline completion

The confirmatory baseline matrix completed in full: 720/720 tuples succeeded with protocol integrity (all 40 turns, no early stopping).

4.2 Condition-level collapse patterns

Table 1: Condition-level collapse summary (N=180 each).

Condition	Mean	Median	SD	N
HOMO_A (Llama)	0.457	0.475	0.311	180
HOMO_B (Qwen)	0.773	0.850	0.210	180
HOMO_C (Mistral)	0.141	0.025	0.208	180
HETERO_ROT	0.250	0.138	0.281	180

4.3 Detector reliability

The preregistered reliability criterion ($\kappa \geq 0.80$) was not met. Although raw agreement was high (92.8%), estimated reliability remained at $\kappa = 0.566$ under high collapse prevalence in the audit set.

5 Discussion

Under fixed protocol settings, collapse dynamics varied substantially by condition. The observed ordering (HOMO_B > HOMO_A > HETERO_ROT > HOMO_C) was stable in this dataset and provides a useful descriptive baseline for follow-on work.




Figure 1: Mean collapse rate by condition with standard deviation error bars.

Path B constraints are central to interpretation. Because the preregistered reliability gate was not met, these results should be used as detector-defined descriptive evidence only. They support prioritizing stronger human-labeled reliability procedures and alternative detector formulations before any confirmatory or causal claim expansion.

6 Limitations

The preregistered detector reliability gate was not met ($\kappa = 0.566$; threshold 0.80). Three rubric iterations were conducted (v1.0: $\kappa = 0.011$; v2.1: $\kappa = 0.519$; v3.0: $\kappa = 0.566$). The gap between raw agreement (92.8%) and κ is consistent with strong prevalence skew. A prevalence-adjusted statistic (PABAK = 0.856) is reported as supplementary sensitivity information, not as gate replacement.

Consequently, all findings are descriptive and condition-comparative rather than detector-validated.

7 Conclusion

We provide a fully closed confirmatory baseline (720/720 tuples) for condition-dependent collapse behavior in 7B-model self-play under a locked protocol. The present contribution is a reproducible descriptive benchmark under Path B claim governance; detector-validation and intervention efficacy remain open tasks for future work.




Figure 2: Collapse rate distribution by condition (boxplot with jittered strip overlay).

References

- Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *International Conference on Learning Representations*, 2020. URL <https://arxiv.org/abs/1904.09751>.
- Jiwei Li, Will Monroe, and Dan Jurafsky. A diversity-promoting objective function for neural conversation models. In *Proceedings of NAACL-HLT*, pages 110–119, 2016. doi: 10.18653/v1/N16-1014.
- Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, pages 3982–3992, 2019. doi: 10.18653/v1/D19-1410.
- Sean Welleck, Ilia Kulikov, Stephen Roller, Emily Dinan, Kyunghyun Cho, and Jason Weston. Neural text generation with unlikelihood training. In *International Conference on Learning Representations*, 2020. URL <https://arxiv.org/abs/1908.04319>.